



559

Non-linear Minimization in R

Fish 559; Lecture 89

Non-linear Minimization Functions in R

559

- Single-dimension methods
 - **optimize** – find the minimum (or maximum) of a univariate function within pre-specified bounds.

Non-linear Minimization Functions in R (Multi-dimension methods)

559

- **nls** – fit a regression model using non-linear least squares. This function produces almost all of the diagnostics associated with **lm**, **lme** and **nlme**.
- **dfp** – minimize an objective function using the Davidson-Fletcher-Powell algorithm and compute 95% intervals (in library **Bhat**).
- **mle** - minimize a real-valued function f subject to constraints (in library **stats4**). Assumes fx in `nll`, computes hessian and inverts it for you – gives you variance
- **newton** – minimize an objective function using a Newton-Raphson algorithm and compute 95% intervals (in library **Bhat**).
- **optim** – minimize a real-valued function f subject to constraints. Family of functions `optimx()` – calls multiple minimization methods for you.
- **nlminb** – minimizes an constrained (or unconstrained) function.

Non-linear Minimization in R

- **Optim/mle** applies one of several methods (quasi-Newton [BFGS], simplex [Nelder-Mead], and conjugate-gradient[CG]) and can make use of derivatives and the Hessian matrix (if they are available). It also includes options to use Simulated Annealing (SANN) – monte-carlo method, very robust but very slow).

Calling **optim-I**

optim(pars, fn, gr, method, control, lower, upper, hessian)

- Required parameters:

- **pars** – a vector of parameters (the initial guesses).
- **fn** – the objective function – that which is to be minimized.

- Optional parameters:

- **gr** – function that computes the gradient.
- **hessian** – should the Hessian matrix be computed?

Calling **optim-II**

- Optional parameters:
 - **method** – which of the various types of minimization algorithm to use.
 - **control** – parameters used to control the minimization (maximum iterations, function calls, etc.).
 - **rel.tol** – how much change in the function before you stop 10e-10 means it's going to try very hard to get a minimum
 - **lower / upper** – vector of lower and upper bounds for the parameters (including $-\text{Inf}$ / Inf).

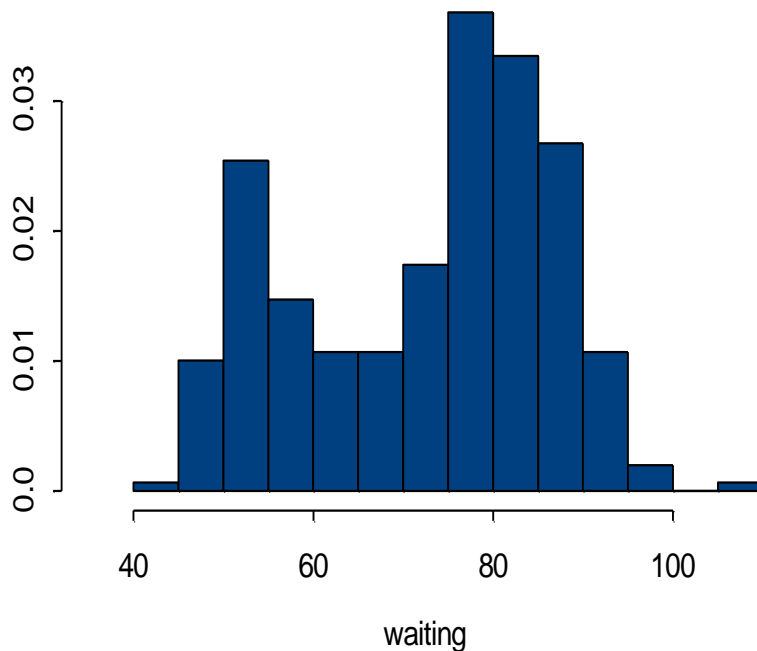
Calling **optim**-III

- (Key) outputs:
 - **par** – the estimates of the parameters.
 - **value** – the lowest value of the objective function.
 - **counts** – number of function and gradient calls.
 - **convergence** – code indicating whether convergence occurred. 0 means it thinks it converged >1 means not
 - **message** – additional information
 - **hessian** – estimate of the Hessian matrix at the solution.
 - **Deviance** – twice the nll making 3.84 the critical value, should be 1.92 for likelihood profile. 1.96 is $2 \times \text{se}$ assuming asymptotic normal.

A first example-I (The waiting data)

559

- The times between eruptions of “Old Faithfull” (dataset “waiting”) appear to be bimodal. We would like to fit these data to a model that is a mixture of two normals.



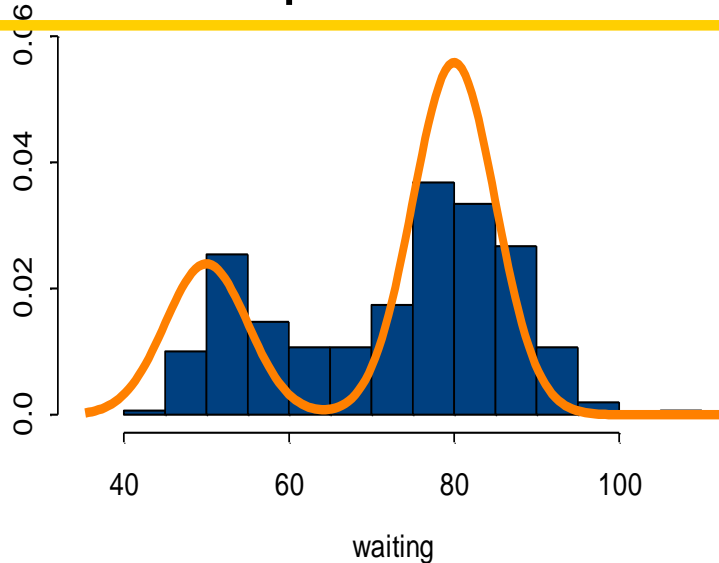
$$P(x) = \pi N(x | \mu_1, \sigma_1) + (1 - \pi) N(x | \mu_2, \sigma_2)$$

We can get sensible estimates for the parameters of this model (0.3, 50, 5, 80, 5) from the histogram.

```
truehist(geyser$waiting,xlim=c(35,110),h=5)
```


A first example-II (The waiting data)

The initial guesses for the parameters are not bad but we can improve on them.



We first write a function that computes the negative log-likelihood.

$$\sum_{i=1}^n \ell n \left[\frac{\pi}{\sigma_1} \phi \left(\frac{y_i - \mu_1}{\sigma_1} \right) + \frac{1-\pi}{\sigma_2} \phi \left(\frac{y_i - \mu_2}{\sigma_2} \right) \right]$$

```
mix.obj <- function(pp,x) {  
  e <- pp[1]*dnorm((x-pp[2])/pp[3])/pp[3]+  
    (1-pp[1])*dnorm((x-pp[4])/pp[5])/pp[5]  
  return(-sum(log(e))) }
```

A first example-III (The waiting data)

559

We can now use **optim** to minimize the negative log-likelihood function.

Initial values

```
wait.init <- c(0.3,50,5,80,5)
```

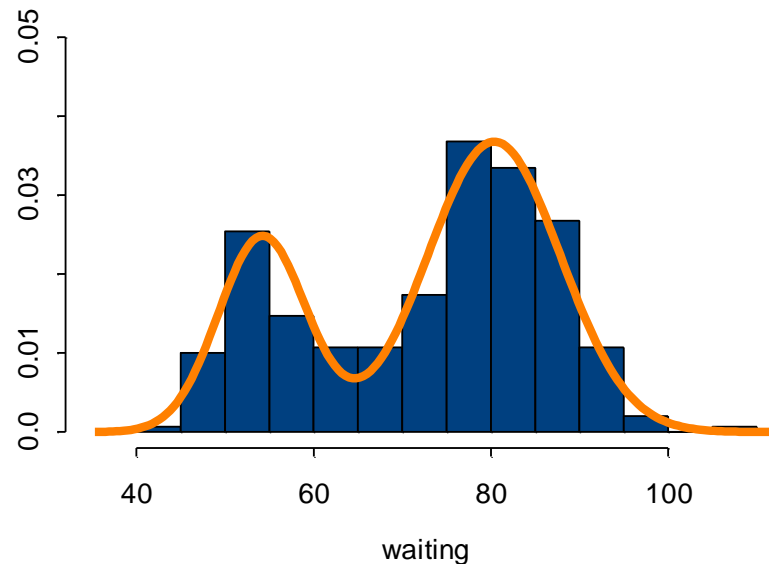
```
mix.n1 <- optim(wait.init, mix.obj, method='L-BFGS-B',  
lower=c(0,-Inf,0,-Inf,0), upper=c(1,rep(Inf,4)), x=geyser$waiting)
```

Place bounds on the
parameter values

Pass the data

A first example-IV (The waiting data)

- The results object should be examined to identify that the minimum has indeed been found.



A first example-V (The waiting data)

559

Our analysis did not use information on the derivatives of the likelihood function. This is a case where it is possible (and relatively straightforward) to make use of derivative information.

```
lmix2a <- deriv( ~log(p*dnorm((x-u1)/s1)/s1 + (1-p)*dnorm((x-u2)/s2)/s2),  
  c("p","u1","s1","u2","s2"),  function(x,p,u1,s1,u2,s2) NULL)
```

```
mix.gr <- function(p,x) {  
  u1 <- p[2]; s1 <- p[3]; u2 <- p[4]; s2 <- p[5]; p <- p[1]  
  e <- lmix2a(x,p,u1,s1,u2,s2)  
  rep(1,length(x)) %*% attr(e,"gradient") }
```

```
mix.n11 <-optim(wait.init, mix.obj, mix.gr, method="L-BFGS-B",  
  lower=c(0,-Inf,0,-Inf,0),upper=c(1,Inf,Inf,Inf,Inf),x=waiting)
```

Calling **mle** - I

mle(minuslogl, start, method, fixed, lower, upper)

- **minuslogl** – the negative log-likelihood function – that which is to be minimized.
- **start** – a named list of estimated parameters (the initial guesses).
- **fixed** – a named list of parameters which are passed to the objective function.
- **method** – which of the various types of minimization algorithm to use.

Calling **mle** - II

- The methods associated with **mle** can be used to:
 - Extract the log-likelihood (**logLik**)
 - Extract the estimated parameters (**coef**)
 - Extract the variance-covariance matrix (**vcov**)
- **mle** provides estimates of the standard errors of the estimates automatically (unlike **optim**, which require you invert the Hessian matrix).
- Hint: **mle** makes computing profile likelihood very easy (compared to **optim**)

The Example Problem

- Fit the dynamic Schaefer model to the catch and effort data for Cape hake off the west coast of South Africa.

$$B_{t+1} = B_t + rB_t(1 - B_t / K) - C_t; \quad B_{1917} = K$$

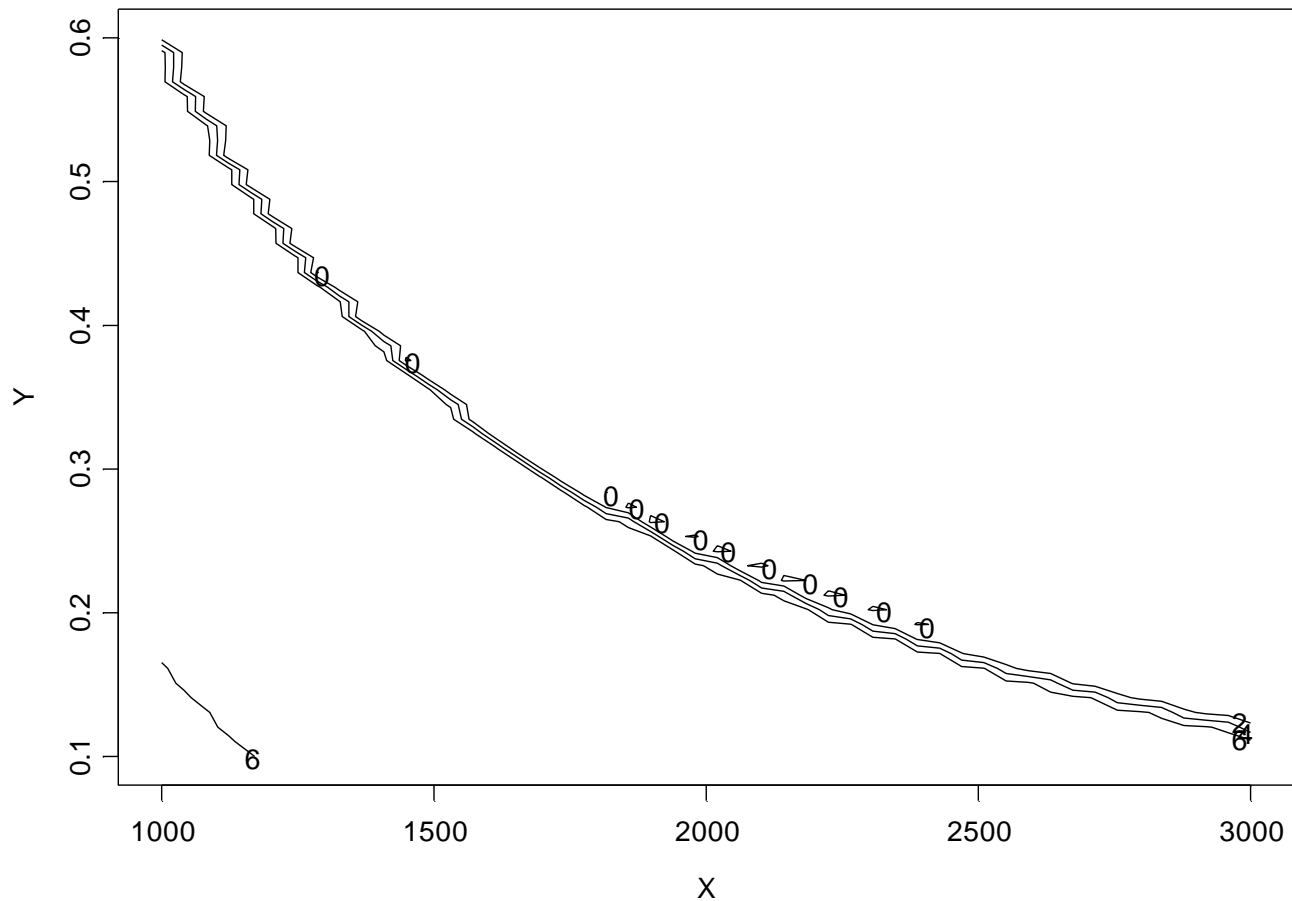
$$-\ell n L = \sum_y \left(\ell n I_y - \ell n(q B_t) \right)^2$$

- We will make use of the maximum likelihood estimate for q :

$$\ell n q = \frac{1}{n} \sum_y \ell n(I_y / B_y)$$

The Example Problem (The sum of squares surface)

559



The Example Problem (The methods)

559

- We could apply methods that are:
 - Derivative free:
 - Simplex / Powell's method
 - Derivative based:
 - `optim`
- Note: that we will not supply **`optim`** with analytic derivatives nor use **`deriv`** to compute them for us (why?)

Passing Parameters

- The likelihood function depends on data (catches, CPUE) as well as on the values for the model parameters (r , q and K).
- It is necessary therefore to pass the data to the function being minimized.
- There are several ways to do this but I prefer to place common parameters in frame 1 so that my function can "see" them.

```
co <- list()
co$Final <- F
co$Nyear <- Nyear
co$Catch <- TheData$Catch
co$CPUEObs <- TheData$CPUE
assign("co",co,pos=1)
```

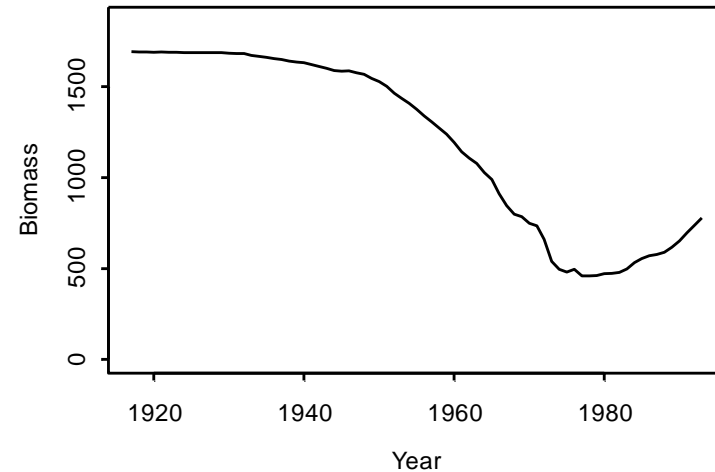
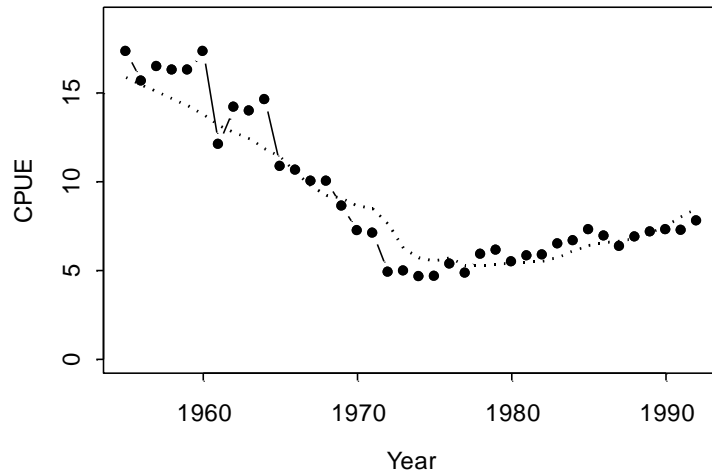
559

The Function to Minimize

```
f1 <- function(x)
{
  K <- exp(x[1]); r <- exp(x[2])
  CpuePred <- rep(0,co$Nyear); Biomass <- rep(0,co$Nyear+1)
  # Do projections
  Biomass[1] <- K
  for (Year in 1:co$Nyear)
    { Biomass[Year+1] <- Biomass[Year] + r*Biomass[Year]*(1.0-Biomass[Year]/K) -
      co$Catch[Year]
    if (Biomass[Year+1] < 0.01) Biomass[Year+1] <- 0.01  }
  # Calculate the ML estimate of q
  qbar <- 0; npar <- 0
  for (Year in 1:co$Nyear)
    if (co$CPUE[Year] > 0)
      { npar <- npar + 1; qbar <- qbar + log(co$CPUE[Year]/Biomass[Year]) }
  qbar <- exp(qbar / npar)
  # Find the SS
  SS <- 0
  for (Year in 1:co$Nyear)
    if (co$CPUE[Year] > 0)
      { CpuePred[Year] <- qbar*Biomass[Year]
        SS <- SS + log(CpuePred[Year]/co$CPUE[Year])^2  }
  return(SS) }
```

559

Results - optim



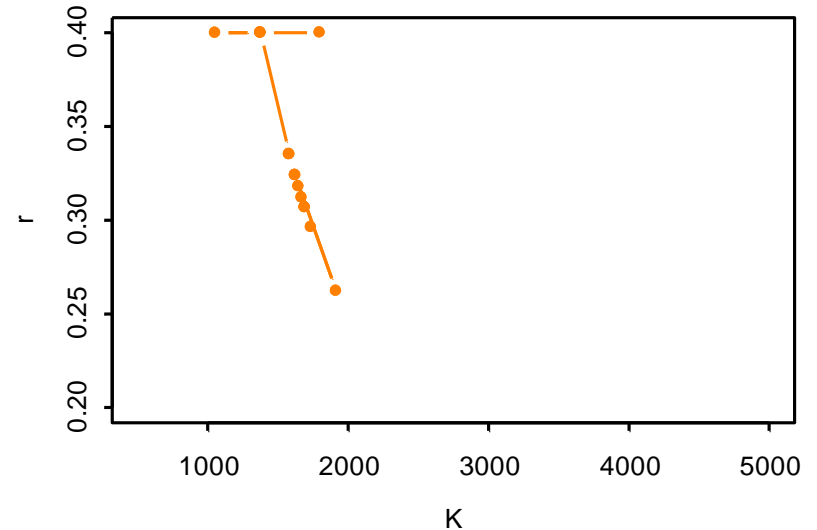
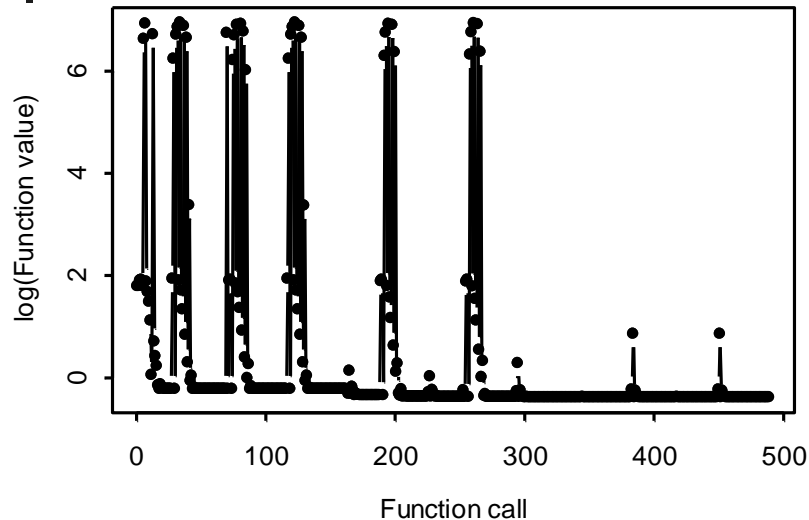
SS=0.6924
 $r=0.335$
 $K=1562$

Function evaluations:65
Gradient evaluations:65

Note: the graphs were produced
by adding code to the function.
The graphs are output if 'final'=T

559

Results - Powell



SS=0.6798

$r=0.307$

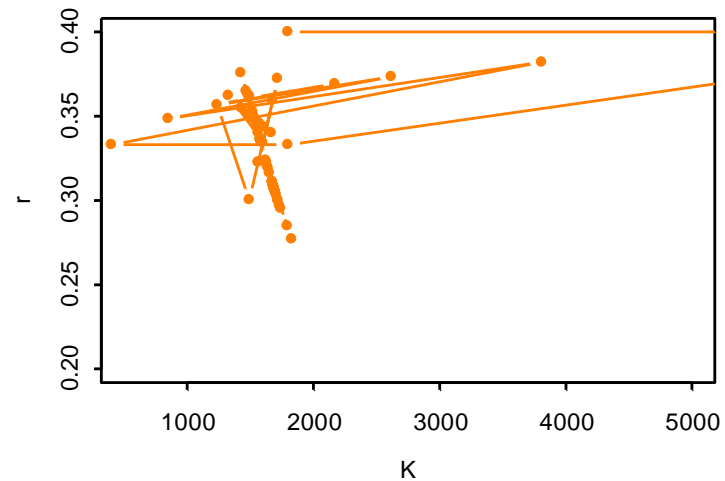
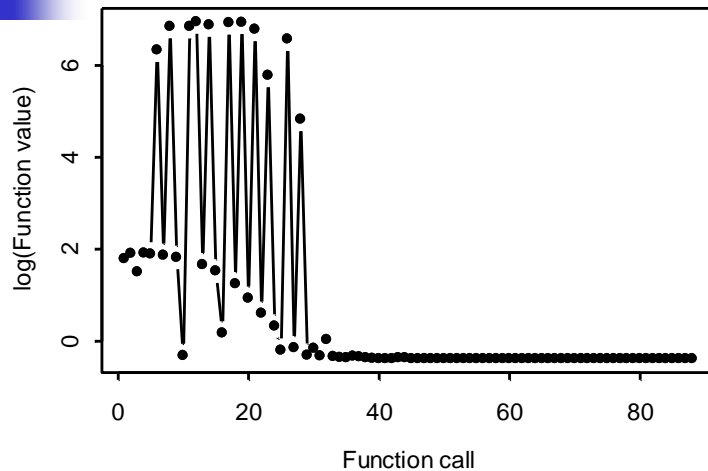
$K=1692$

468 function evaluations!

The algorithm requires some tuning!

559

Results - Simplex



$SS=0.6798$

$r=0.307$

$K=1692$

88 function evaluations

45 iterations

Comparison of Packages (from Schnute *et al.*, 1998)

559

Parameters	Product	Calls	Times
37	ADMB	161	4.6 s
	Gauss	4041	2.8 min
	MatLab	1936	5.8 min
	S-plus	n/a	n/a
100	ADMB	291	38 s
	Gauss	23,365	1.08 hr
	MatLab	18,360	3.25 hr
	S-plus	n/a	n/a

The trials involved fitting an age-structured model₂₃

Hints

- Methods may work better if the variables are all scaled to approximately 1.
- If some of the parameters are linear functions of the parameters (e.g. the q and σ parameters in a likelihood function for the Schaefer model), find analytical solutions for them – this reduces the dimension of the problem.
- Never trust a non-linear minimizer!
 - Try different starting values.
 - Combine different methods (e.g. apply Simplex and then a conjugate gradient method). The function `optimx` in the package `optimx` provides a wrapper for the algorithms in `optim` (except `SANN`) as well in other packages.
 - Test your code with deterministic data, then simulated data with stochastic processes.